Instructions for EMNLP-IJCNLP 2019 Proceedings

Anonymous EMNLP-IJCNLP submission

Abstract

1 Introduction

Expressions describing or referring to objects in visual scenes typically include a word naming the type of the object: e.g., *cheesecake* or *dessert* in Figure **??**. Determining these objects names is a core aspect of virtually every language & vision task, ranging from e.g. referring expression generation to visual dialogue. Nevertheless, research in language & vision has mostly sidestepped questions about how speakers actually choose these names and how computational models should account for it.

While state-of-the-art computer vision systems are able to accurately classify images into thousands of different categories (e.g. **?**), they mostly adopt very simple assumptions with respect to the underlying lexicon, which is typically implemented as a simple, flat labeling scheme. Thus, a standard object recognition system would be trained to classify the objects in Figure 1 as either *dessert* or *cake*. In contrast, humans seem to be more flexible as to the chosen level of generality and to the chosen part of the taxonomy (see objects in Figure 1 that could be named *cake*, *cheesecake*, *dessert*, *sweet*, *pastry*, *food* etc.) Seminal work on prototypes suggests that the prototypicality of the object will determine the level of generality of the object name, i.e. a robin can be named *bird*, but a penguin is better referred to as "*penguin*" (**?**).

Two main findings in the literature:

- very high agreement, most people use the same name for the same object
- prototypicality is a factor, context too, but research has only looked at generality/specificity of the name

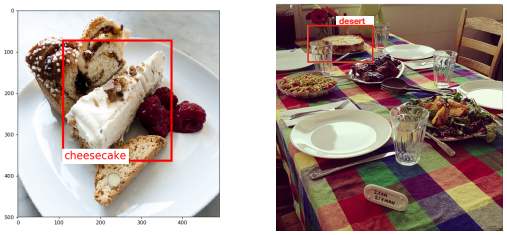

Figure 1: Two objects of the same type of cake, with different names in VisualGenome

//sz: something is missing here .. explain why exactly we did what we did, why is it interesting to collect many names for the same object?//

There are two main findings:

- the level of agreement in object naming is much higher in certain domains than in others, as it happens, the domains that have been traditionally used in object naming research (e.g. animals) seem to display the highest amount of agreement in our data set
- while previous work has mostly focussed on variation in the level of generality (*penguin* vs. *bird*), our datasets contains a lot of variability for names coming from different parts of the taxonomy (*dessert* vs. *cake*, *bottle* vs. *wine*)

2 Related Work

3 Analysis

3.1 Agreement

We compute the following agreement measures:

- **top %**: for each object, we calculate the relative frequency of the most common name, and then average over all objects
- **SD %**: for each object, we calculate the Snodgrass agreement measure, and then average over all objects

- **=VG**: the proportion of objects where the most frequent name coincides with the name annotated in VisualGenome

Table 1 shows that, overall, our annotators achieve a fair amount of agreement in the object naming choices. The domain where annotators agree most is the animal domain, which, interestingly, happens to be the domain that has been mostly discussed in the object naming literature. *//sz: ... much more to say//*

Why is naming more flexible in certain domains than in others?

3.2 Lexical relations

In this section, we take a closer look at the lexical variation we observe in our data set. We analyze the data points where participants attributed different names to the same object and extract a set of pairwise **naming variants**. These naming variants correspond to pairs of words that can be used interchangeably to name certain objects. For each object, we extract the set of naming variants $s = \{(w_{top}, w_2), (w_{top}, w_3), (w_{top}, w_4), ...\}$ where w_{top} is the most frequent name annotated for the object and $w_2 ... w_n$ constitute the less frequent alternatives of w_{top}. The **type frequency** of a naming variant (w_{top}, w_x) corresponds to the number of objects where this variant occurs. The **token frequency** of (w_{top}, w_x) corresponds the count of all annotations where w_x has been used instead of w_{top}. In Table 3, we show the the naming variants with the highest raw token frequency for each domain. *//sz: domains need to be updated//*

The naming variants can be grouped according to their lexical relation, as follows:

- **synonymy**: e.g. aircraft vs. airplane
- **hyponymy**: e.g. man vs. person
- **co-hyponymy**: e.g. chicken vs. dinner, swan vs. goose
- **no relation**: e.g. desk vs. apple

Research on object naming following the idea of entry-level categories has, essentially, exclusively looked at names that stand in a hierarchical relation (i.e. hyponymy/hypernymy).

We use WordNet to extract lexical relations between the naming variants in our data set. Unfortunately, this means that we have to exclude a certain portion of the data as either (i) one of the name is not covered in WordNet, (ii) we cannot find a lexical relation between the two names (see below). Also, we had to be relatively permissive with respect to the definition of hyponymy/co-hyponymy. For instance, to analyze *giraffe* as a hyponym of *animal* we have to look at the closure of the hyponyms of *animal* with a depth of 8 (in WordNet). *//sz: should we call this co-hyponymy or co-hierarchical relation?//*

//sz: include Table that reports counts of the naming variants, coverage in WordNet etc.//

Table 2 shows the distribution of lexical relations for those naming variants that we were able to analyze with WordNet. Both in terms of their types and token frequency, the naming variants that instantiate a (loose) co-hyponymy relation are by far the most frequent. *//sz: discuss in more detail, discuss: to what extent is this an artefact of WordNet?//* This is really interesting: most research on object naming, to date, has focussed on hyponymy/hypernymy, i.e. variation that relates to hierarchical relations between object names. Our data suggests that co-hierarchical variation is really important too.

3.3 Issues with WordNet

Some (interesting, somewhat cherry-picked) word pairs were WordNet does not find any relation (excluded in the above analysis):

- lettuce – salad
- fruit – food
- man – catcher
- bowl –chili
- bowl – diner
- burger – meat
- statue – animal (image shows statue of an animal)
- bottle – alcohol
- donut –desert
- zebra – stripes
- oven – grill

//sz: discuss...//

	all synsets			max synset				min synset			
domain	% top	SD	=VG	id	% top	SD	=VG	id	% top	SD	=VG
people	0.52	2.13	0.50	professional.n.01	0.61	2.02	0.20	athlete.n.01	0.36	2.62	0.37
clothing	0.64	1.58	0.70	neckwear.n.01	0.79	0.91	0.77	footwear.n.01	0.47	2.55	0.40
home	0.66	1.50	0.78	tool.n.01	0.86	0.73	0.94	crockery.n.01	0.52	1.92	0.40
buildings	0.67	1.55	0.73	bridge.n.01	0.75	1.21	0.87	place_of_worship.n.01	0.46	2.26	0.08
food	0.71	1.30	0.63	edible_fruit.n.01	0.80	0.89	0.79	vegetable.n.01	0.53	1.97	0.15
vehicles	0.72	1.13	0.71	train.n.01	0.93	0.42	0.99	aircraft.n.01	0.52	1.50	0.41
animals,plants	0.91	0.44	0.94	feline.n.01	0.95	0.29	0.99	fish.n.01	0.39	2.53	0.55
all	0.70	1.34	0.73								

Table 1: Agreement in object names for objects of different domains, if applicable, synsets with maximal and minimal agreement (top %) are shown

relation	% types	% tokens	av. depth
co-hyponymy (closure, max depth=10)	0.889	0.551	3.479
hyponymy (closure, max depth=10)	0.097	0.328	2.204
synonymy	0.015	0.121	1.000

Table 2: Lexical relations between naming variants according to WordNet, for the set of name pairs where both words can be found in WordNet and stand in a *//sz: should we produce this table for the different domains?//*

3.4 Entry-level names and preference orders....

//sz: an interesting example:// In our data set, there are 24 images where *penguin* has been used, so we know that the object is a *penguin*. For 50% of these images, annotators still prefer *bird* as the most common name. According to the theory of entry-level categories, this should not happen. People should always prefer *penguin* over *bird*.

//sz: how can we analyze this quantitatively?//

3.5 Co-hyponyms as names for objects

//sz: analyze and discuss why we find so many co-hyponyms in our data set//

4 Modeling

//sz: I think, the prevalence of co-hyponymy is the most interesting finding. Can we learn to predict whether to co-hyponyms can be used to name the same object?//

category	most frequent naming variants
people	woman – person (3594), man – person (3546), boy – child (3243), woman – girl (2328), girl – child (1985), woman – tennis player (1277), man – player (1273), man – boy (1214), skateboarder – skater (1194), man – t-shirt (1143)
food	pizza – food (1883), sandwich – food (1123), hotdog – food (540), pizza – cheese (457), pizza – plate (430), salad – food (402), sandwich – burger (398), hotdog – sandwich (351), sandwich – bread (318), cake – food (286)
home	couch – sofa (4090), desk – table (3448), carpet – floor (1697), bench – chair (1401), desk – keyboard (1380), counter – table (1201), table – desk (1135), counter – countertop (1101), table – counter (906), rug – carpet (895)
buildings	house – building (1160), building – house (511), bridge – train (326), bridge – overpass (235), house – window (161), house – home (123), tent – canopy (120), building – castle (101), bridge – building (98), bridge – pole (85)
vehicles	airplane – plane (11194), plane – airplane (3829), motorcycle – bike (2624), airplane – jet (1319), boat – ship (1301), truck – car (1095), car – vehicle (874), motorcycle – wheel (861), truck – vehicle (718), truck – wheel (716)
clothing	shirt – t-shirt (2914), jacket – coat (2396), jacket – shirt (1552), jacket – suit (1168), suit – jacket (1029), shirt – jacket (813), shirt – tie (723), shirt – man (487), shirt – dress (462), shirt – sweater (450)
animals_plants	cow – bull (515), sheep – goat (486), cow – animal (445), giraffe – animal (380), bird – parrot (349), sheep – animal (294), sheep – lamb (282), horse – animal (269), cat – animal (237), bird – seagull (231)

Table 3: Most frequent naming variants for each category